

Analysing Anomaly Detection Methods For Time Series Using R



Structure

1. Context and Motivation
2. Anomalies and Outcomes
3. Detection Methods and Analysing them using R

Context and Motivation



Switch energy & save

It's quick and easy to save up to £679/yr*

[Compare your gas & electricity](#)

[Fixed energy plan ending? ▶](#)

[Gas and electricity ▶](#)

[Electricity prices ▶](#)



Broadband, TV, landline

Fast, reliable broadband from £2.50

[Compare broadband deals](#)

[Broadband deals ▶](#)

[Broadband & home phone ▶](#)

[Broadband & TV packages ▶](#)



Mobiles

Latest deals from only £10

Credit cards

Compare credit cards

Car insurance

Quick quote and save today!

Mortgages

Compare mortgage rates

Situations

1. A tv show talks about energy savings and causes a big spike in online traffic
2. A bug is released onto an important page in our funnel, causing a big drop in users completing our energy process
3. A product is taken off our energy tariff price comparison page, causing a drop in users proceeding to the next step.
4. News stories over several days causes an increase in Google traffic to one of our guides pages.

How these changes are often noticed

1. 'Oh, traffic has gone up on the site!'
2. 'Why has this KPI dropped in the last two hours?'
3. 'Hmm, this landing page started getting loads of traffic 3 days ago'

Challenges

- Real time data streams can't be checked constantly in a reliable way by people with many other things to do.
- It is time consuming to check thousands of time series regularly for changes

Anomalies and Outcomes

Anomaly: What is it?

- Anomaly: something that is different from what is expected, or not in agreement with something else.
- Expected ranges and values are defined differently by different methods.
- Our definition of anomaly will change with the methods.

Outcomes

1. Quickly and reliably detect large and important anomalies
2. Eventually and reliably detect less important anomalies
3. To indicate where to look to find out why the anomaly occurred
4. Do the above without many false positives

Detection Methods and Analysing them in R

A Few Approaches

1. Manually setting max/min thresholds based on time
2. Robust Principal Component Analysis
3. Changepoint Detection


Setting Maximum and Minimum Thresholds Based on Time

Setting Max and Min Thresholds

1. Choose a metric you record in real time
2. (Optional) transform into a new metric over a period of recent time.
3. Set upper and lower limits for this metric - if the metric goes outside these bounds, detect an anomaly.
 - Limits can depend on time. If you have seasonal data, then the limits can depend on the time in each season.
 - Send an alert to interrupt key people


Pageviews to Important Page By Minute For Example Week

Pageviews



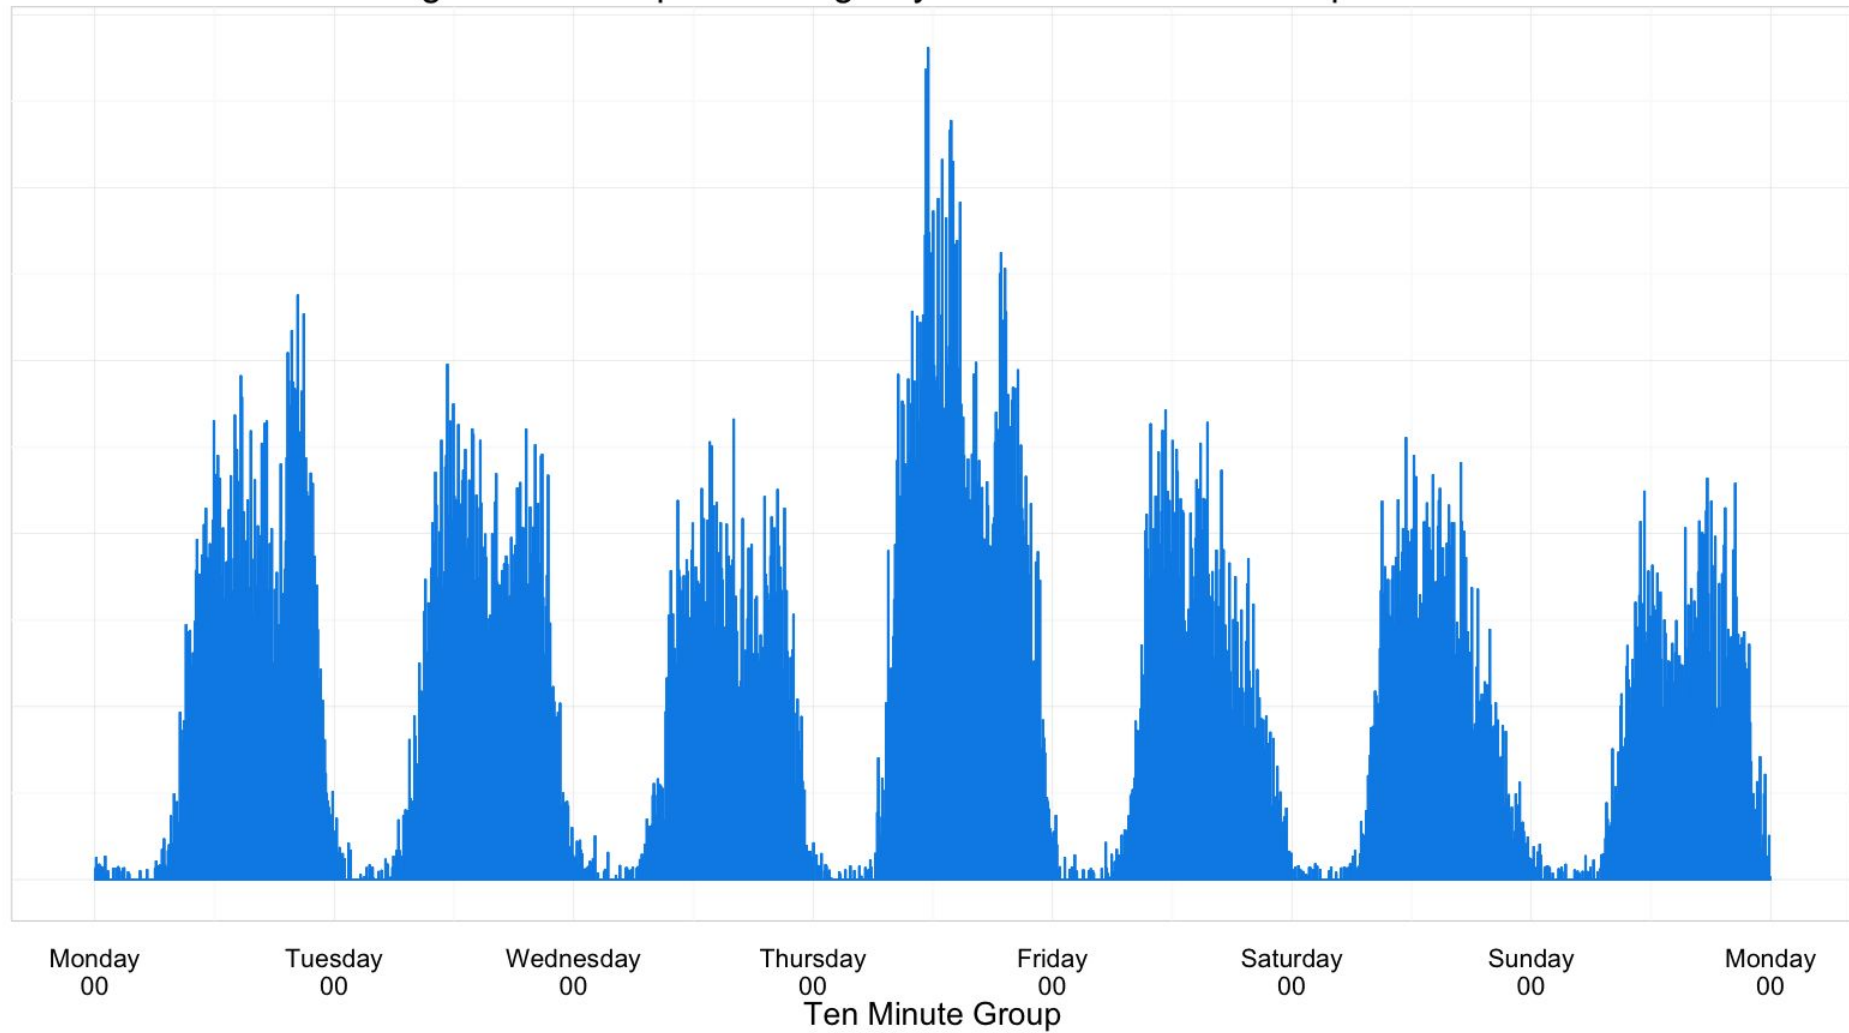
Pageviews to Important Page By Hour For Example Week

Pageviews



Pageviews to Important Page By Ten Minutes For Example Week

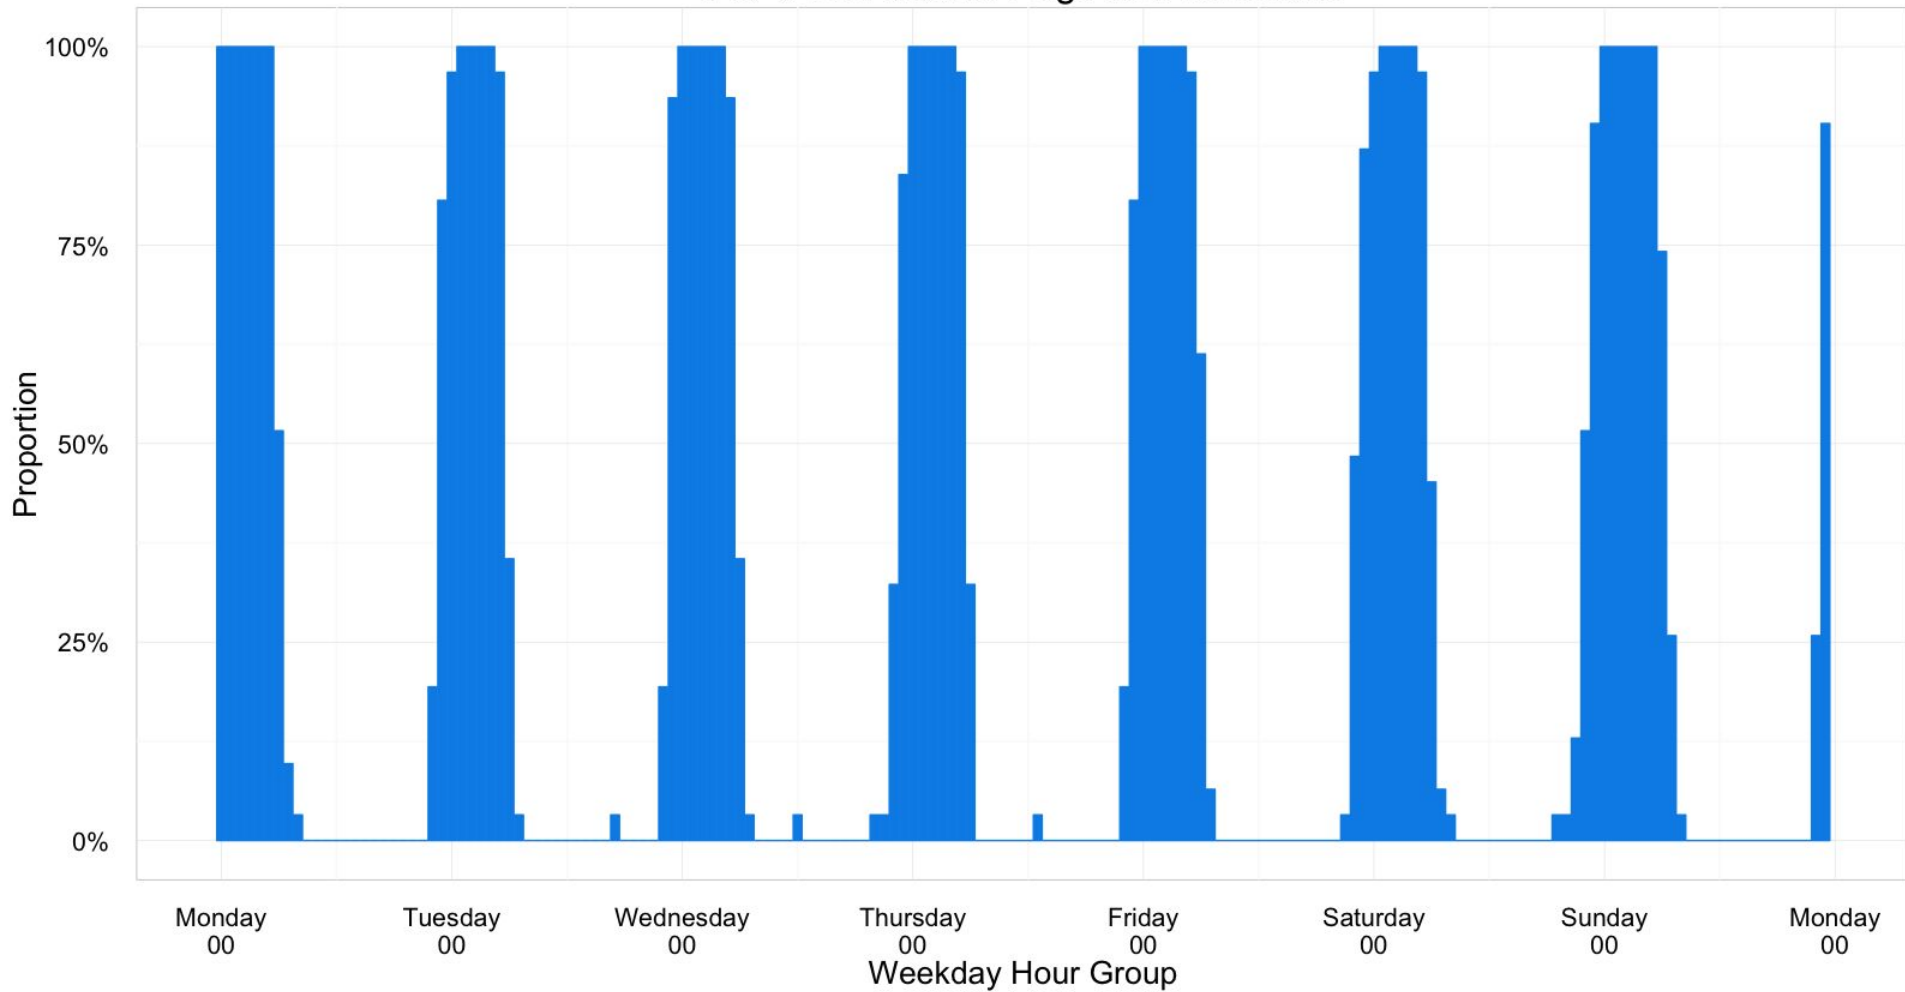
Pageviews



Start Simple

- Extracted 31 weeks of data
- Have ten minute rolling sum buckets
- Create a detector function which classifies a point as anomalous if it is equal to 0
- Group by weekday - hour combination
- For each of these hours, find out what proportion of days in our 31 week dataset would have an alert fired off.


Proportion of Days Anomalies Detected in Weekday Hour
For 0 Ten Minute Pageview Minimum



Next...

- Have a program check the datastream every few minutes outside of those nighttime hours, send an alert if zero pageviews detected.
- We can do better.
- Ignore nighttime hours, try a couple of arbitrary thresholds


average alerts zeromin



average_alerts_highermin



average_alerts_doublehighermin



Monday
00

Tuesday
00

Wednesday
00

Thursday
00

Friday
00

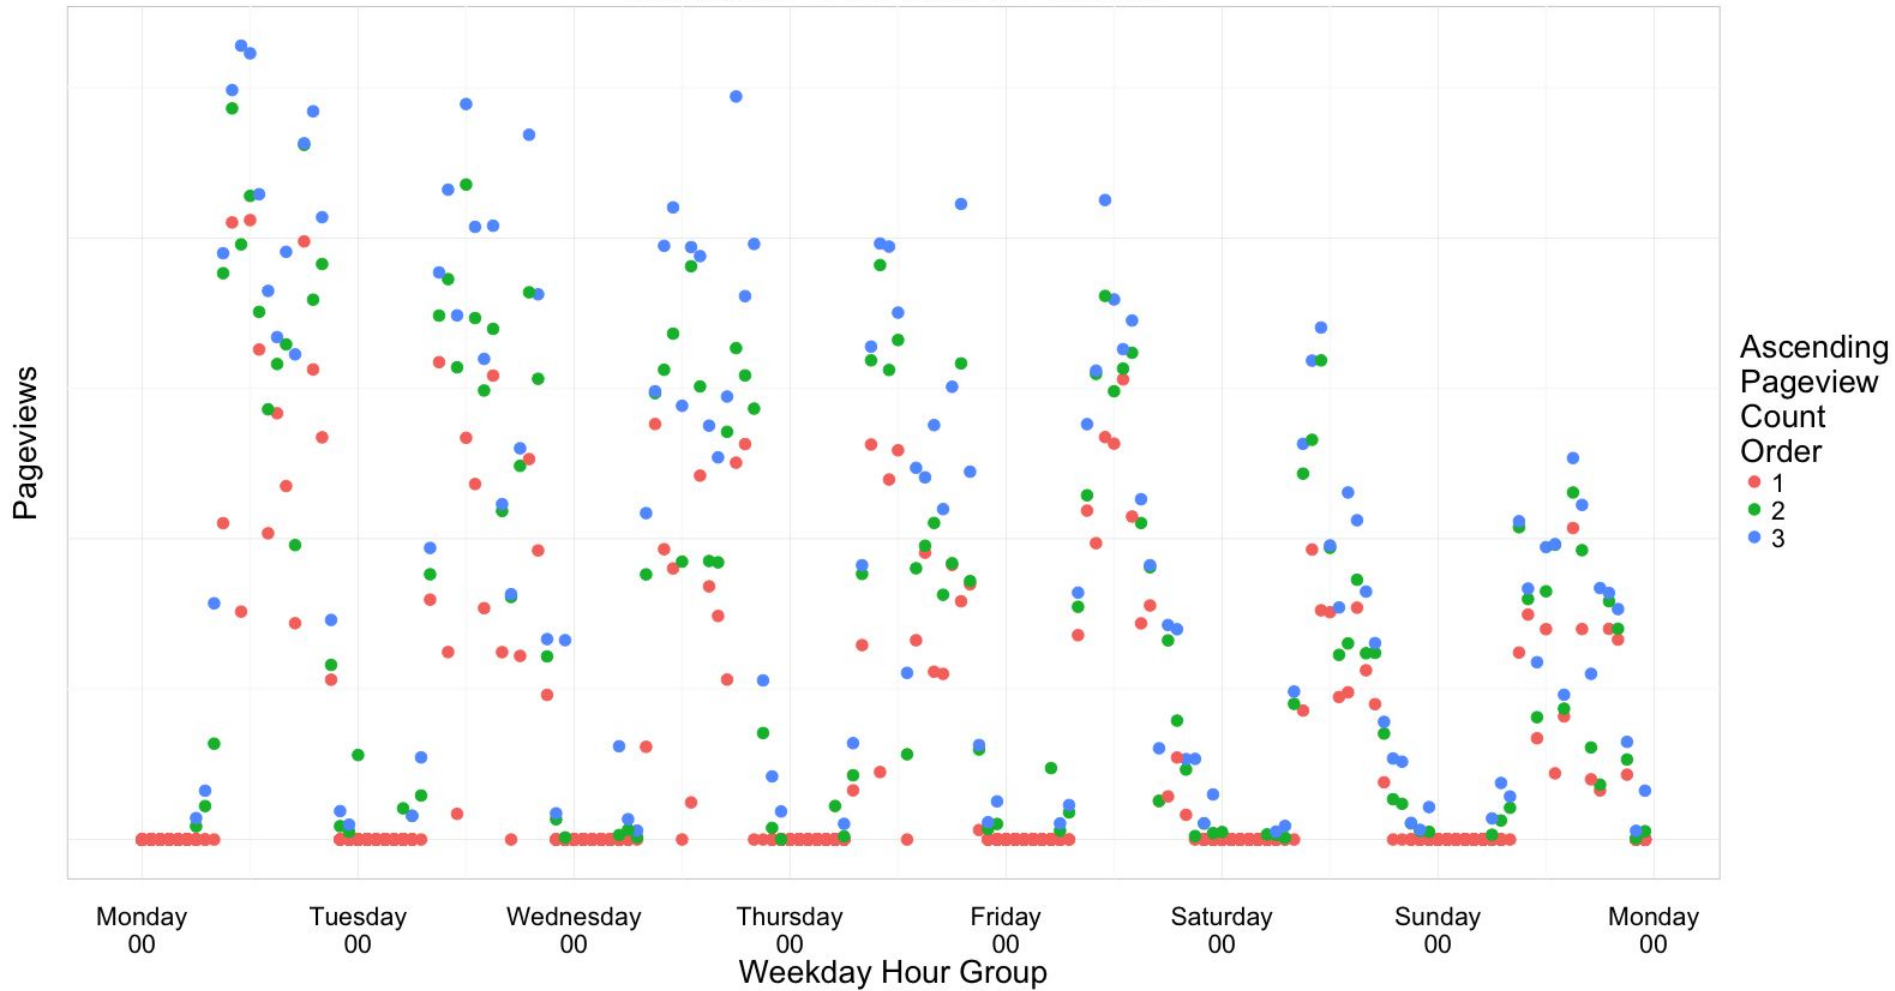
Saturday
00

Sunday
00

Monday
00

Hour Groups


Lowest 3 Ten Minute Important Page Pageviews
Measured Over Last 31 Weeks



Next Step

- Find the minimum pageviews for each weekday hour bucket for each week
- Create 3 detectors:
 - First: detect anomaly if pageviews is less than or equal to the **lowest recorded** value for that weekday hour bucket in the historical dataset
 - Second: detect anomaly if pageviews is less than or equal to the **second lowest** recorded value for that weekday hour bucket in the historical dataset
 - Third: detect anomaly if pageviews is less than or equal to the **third lowest** recorded value for that weekday hour bucket in the historical dataset

Proportion of Days Anomalies Detected by Weekday For Lowest Hour Bucket Minima



Useful Packages

- Dplyr or Data Table
- Lubridate and Zoo
- RcppRoll and Zoo
- Reshape2
- Ggplot2

Robust Principal Component Analysis

Robust Principal Component Analysis


Data Cleaning:

1. Start with a timeseries
2. Choose an interval of time that corresponds to a season in the series
3. Create a matrix D with each row as one of these intervals, in order
4. Whiten each column - mean to zero and variance to 1

Actual RPCA step:

5. Find a way of representing this matrix as a low rank matrix L added to a sparse matrix S added to an error matrix E . Transform into original space.


Is there an anomaly here? Yes. Yes there is.



	t1	t2	t3	t4	t5
s1	2.47	2.43	0.97	1.26	2.63
s2	2.56	2.82	1.64	1.15	1.70
s3	2.92	2.50	24.58	0.94	2.29
s4	2.84	2.90	1.10	0.67	2.97
s5	2.83	2.68	1.60	0.90	2.16



	t1	t2	t3	t4	t5
s1	2.47	2.43	0.97	1.26	2.63
s2	2.56	2.82	1.64	1.15	1.70
s3	2.92	2.50	1.58	0.94	2.29
s4	2.84	2.90	1.10	0.67	2.97
s5	2.83	2.68	1.60	0.90	2.16



	t1	t2	t3	t4	t5
s1	0	0	0	0	0
s2	0	0	0	0	0
s3	0	0	23	0	0
s4	0	0	0	0	0
s5	0	0	0	0	0

Use Netflix Surus Package

```
> library(devtools)
> install_github(repo = "Surus", username = "Netflix", subdir = "resources/R/RAD")
```

```
> ?AnomalyDetection.rpca
```

```
AnomalyDetection.rpca(X, frequency = 7, dates = NULL, autodiff = T,
  forcediff = F, scale = T, L.penalty = 1,
  s.penalty = 1.4/sqrt(max(frequency, ifelse(is.data.frame(X), nrow(X),
length(X))/frequency))), verbose = F)
```

Use Netflix Surus Package

Value


- X_transform. The transformation applied to the time series, can be the identity or could be differencing
- L_transform. The low rank component in the transformed space
- S_transform. The sparse outliers in the transformed space
- E_transform. The noise in the transformed space

Apply to Data


- 31 weeks, 168 hours each week of pageviews
- Use default S penalty
- Prioritise our weeks.
- Rank each week with the maximum absolute S value recorded

Week with Max Absolute S Rank: 1


important_page_pageviews



S_transform



L_transform



Monday
00

Tuesday
00

Wednesday
00

Thursday
00

Friday
00

Saturday
00


Sunday
00

Monday
00

Hour of Week

Week with Max Absolute S Rank: 2

important_page_pageviews




S_transform

Pageviews



L_transform



Monday
00

Tuesday
00

Wednesday
00

Thursday
00

Friday
00

Saturday
00


Sunday
00

Monday
00

Hour of Week


Week with Max Absolute S Rank: 3

important_page_pageviews




S_transform

Pageviews



L_transform



Monday
00

Tuesday
00

Wednesday
00

Thursday
00

Friday
00

Saturday
00

Sunday
00

Monday
00

Hour of Week

```
AnomalyDetection.rpca(X, frequency = 7, dates = NULL, autodiff = T,  
  forcediff = F, scale = T, L.penalty = 1,  
  s.penalty = 1.4/sqrt(max(frequency, ifelse(is.data.frame(X), nrow(X),  
length(X))/frequency))), verbose = F)
```

Function Being Minimized

$0.5 * \text{EuclideanNorm}(X - L - S)^2$

$+ \text{LPenalty} * \text{NuclearNorm}(L)$

$+ \text{SPenalty} * \text{L1Norm}(S)$

Changepoints

Changepoints for Means

1. Choose a timeseries
2. Assume that a known probability distribution with unknown mean generates the data at each point.
3. Find a likely partition on which the mean is constant on each segment.


Changepoint Package

- Available on CRAN
- Can find changepoints for both variances and means
- Use `cpt.mean` for detecting changes in mean.

```
cpt.mean(data,penalty="MBIC",pen.value=0,method="AMOC",Q=5,test.stat="Normal",class=TRUE,  
param.estimates=TRUE,minseglen=1)
```



$y=c(\text{rnorm}(25, 15, 2), \text{rnorm}(15, 20, 2), \text{rnorm}(20, 10, 4))$

Method = "AMOC" Penalty = "MBIC"




$y=c(\text{rnorm}(25, 15, 2), \text{rnorm}(15, 20, 2), \text{rnorm}(20, 10, 4))$


Method = "PELT" Penalty = "MBIC"



$y=c(\text{rnorm}(25, 15, 2), \text{rnorm}(15, 20, 2), \text{rnorm}(20, 10, 4))$
Method = "PELT" Penalty = "Manual" pen.value = 200



$y=c(\text{rnorm}(25, 15, 2), \text{rnorm}(15, 20, 2), \text{rnorm}(20, 10, 4))$
Method = "PELT" Penalty = "Manual" pen.value = 100



Cpt.mean outputs vectors with length = no. of changepoints.

```
cpt.mean.vec <- function(x, method = 'PELT', penalty = 'MBIC', v) {  
  m <- cpt.mean(x, method = method, penalty = penalty, pen.value = v)  
  m_cpts <- m@cpts  
  reps <- m@cpts - lag(m@cpts, default = 0)  
  m_next_cpts <- rep(m_cpts, reps)  
  m_means <- m@param.est$mean  
  m_next_step_size <- rep(c(m_means[2:length(m_means)] - m_means[1:(length(m_means)-1)]), 0), reps)  
  m_current_means <- rep(m_means, reps)  
  output <- data.frame(current_means = m_current_means, next_cpt = m_next_cpts, next_step_size = m_next_step_size)  
  return(output)  
}
```

Useful Working Pattern

```
> head(df, 15)
```

	segment	timestamp	group	metric
1	1	1	A	23.7
2	1	2	A	22.0
3	1	3	A	18.2
4	1	4	A	19.3
5	1	5	A	17.8
6	2	6	A	14.9
7	2	7	A	21.9
8	2	8	A	26.3
9	2	9	A	26.1
10	2	10	A	27.8
11	1	1	B	21.3
12	1	2	B	28.6
13	1	3	B	37.0
14	1	4	B	28.0
15	1	5	B	26.3

Useful Working Pattern

```
> dfdm <- cbind(df, detection_metrics); head(dfdm, 15);
```

	segment	timestamp	group	metric	detection_metric1	detection_metric2
1	1	1	A	23.7	21.3	0.073
2	1	2	A	22.0	21.3	0.069
3	1	3	A	18.2	21.3	0.106
4	1	4	A	19.3	21.3	0.002
5	1	5	A	17.8	21.3	0.191
6	2	6	A	14.9	21.3	0.106
7	2	7	A	21.9	21.3	0.165
8	2	8	A	26.3	21.3	0.004
9	2	9	A	26.1	21.3	0.135
10	2	10	A	27.8	21.3	0.105
11	1	1	B	21.3	24.7	0.236
12	1	2	B	28.6	24.7	0.339
13	1	3	B	37.0	24.7	0.213
14	1	4	B	28.0	24.7	0.357
15	1	5	B	26.3	24.7	0.321

Useful Working Pattern

```
> dfdm <- group_by(dfdm, segment, group) %>%  
+   mutate(max_seggroup_detect2_value = max(detection_metric2)) %>%  
+   ungroup %>%  
+   mutate(group_segment_ranking = dense_rank(-max_seggroup_detect2_value))
```


Useful Working Pattern

```
> head(dfm, 15)
```

```
# A tibble: 15 × 9
```

	segment	timestamp	group	metric	detection_metric1	detection_metric2	max_seggroup_detect2_value	group_segment_ranking	max_detect2_value
	<int>	<int>	<fctr>	<dbl>	<dbl>	<dbl>	<dbl>	<int>	<dbl>
1	1	1	A	23.7	21.3	0.073	0.191	5	0.191
2	1	2	A	22.0	21.3	0.069	0.191	5	0.191
3	1	3	A	18.2	21.3	0.106	0.191	5	0.191
4	1	4	A	19.3	21.3	0.002	0.191	5	0.191
5	1	5	A	17.8	21.3	0.191	0.191	5	0.191
6	2	6	A	14.9	21.3	0.106	0.165	6	0.165
7	2	7	A	21.9	21.3	0.165	0.165	6	0.165
8	2	8	A	26.3	21.3	0.004	0.165	6	0.165
9	2	9	A	26.1	21.3	0.135	0.165	6	0.165
10	2	10	A	27.8	21.3	0.105	0.165	6	0.165
11	1	1	B	21.3	24.7	0.236	0.357	4	0.357
12	1	2	B	28.6	24.7	0.339	0.357	4	0.357
13	1	3	B	37.0	24.7	0.213	0.357	4	0.357
14	1	4	B	28.0	24.7	0.357	0.357	4	0.357
15	1	5	B	26.3	24.7	0.321	0.357	4	0.357


Useful Working Pattern

```
> dfdm.melt <- melt(dfdm, c('segment', 'timestamp', 'group', 'group_segment_ranking'),  
+                   measure.vars = c('metric', 'detection_metric1', 'detection_metric2'))  
> head(dfdm.melt, 15)
```

	segment	timestamp	group	group_segment_ranking	variable	value
1	1	1	A	5	metric	23.7
2	1	2	A	5	metric	22.0
3	1	3	A	5	metric	18.2
4	1	4	A	5	metric	19.3
5	1	5	A	5	metric	17.8
6	2	6	A	6	metric	14.9
7	2	7	A	6	metric	21.9
8	2	8	A	6	metric	26.3
9	2	9	A	6	metric	26.1
10	2	10	A	6	metric	27.8
11	1	1	B	4	metric	21.3
12	1	2	B	4	metric	28.6
13	1	3	B	4	metric	37.0
14	1	4	B	4	metric	28.0
15	1	5	B	4	metric	26.3


Useful Working Pattern

```
> filter(dfdm.melt, group_segment_ranking == 1) %>%  
+   ggplot(aes(timestamp, value, fill = variable, group = variable)) +  
+   geom_bar(stat = 'identity') +  
+   facet_grid(variable~group, scale = 'free') +  
+   guides(fill = FALSE) +  
+   theme(text = element_text(size = 20))
```




Week with Max Absolute S Rank: 1


important_page_pageviews



S_transform



L_transform



Monday
00

Tuesday
00

Wednesday
00

Thursday
00

Friday
00

Saturday
00

Sunday
00

Monday
00

Hour of Week

Summary - Real Time

- Do it. Start simple and iterate.
- Investigate boring but reliable methods first.
- Use historical data to figure out how often alerts would be fired off

Summary - Outside of Real-Time

- Prioritise, don't classify
- RPCA (RAD package) for finding spikes and dips
- Changepoints (changepoint package) for finding longer-term changes
- Adjust penalty values to control sensitivity.
- Rank time series with understandable metrics first based on:
 - Changepoint mean differences
 - S values
 - Different penalty values

Thank You!

shannon.wirtz@uswitch.com

We're recruiting: <https://www.uswitch.com/vacancies/>